

Deploy a Full-Stack App with AI Coding Assistants

Part 3 of AI Dev Tools Zoomcamp



ALEXEY GRIGOREV

AUG 10, 2026



19



3

Share

This is the third article in a series for [AI Dev Tools Zoomcamp](#), the free course we run at DataTalks.Club.

All articles in the series:

- Part 1: [AI-Native Development: Specifications, Loop and Graph Engineering](#)
- Part 2: [Build and Ship a Full-Stack App with AI Coding Assistants](#)
- Part 3: [Deploy a Full-Stack App with AI Coding Assistants](#) (this article)
- Part 4: [DevOps and Observability for an AI-Built App](#)
- Part 5: [Coding Agent Building Blocks: Reusable Skills and Specialized Subagents](#)

Thanks for reading Alexey On Data! Subscribe for free to receive new posts and support my work.



Test, Containerize, and Deploy an AI-Assisted App - Alexey Grigorev

DataTalksClub



Regarder sur

In this article, I want to take a web application and make it accessible for everyone on the internet.

We will:

- Take the application that's already working locally
- Containerize the application
- Switch from SQLite to Postgres
- Add integration tests
- Deploy to AWS
- Set up CI/CD via GitHub Actions

It's based on the second half of the full-day workshop I did for AI Shipping Labs: [Build and Deploy a Full-Stack App with AI Coding Assistants](#).

Recap

In the [previous article](#), we started building an application for system-design interviews. An interviewer creates a session and shares a link with a candidate. When the candidate updates something on the canvas, the interviewer sees the updates in real time.

- First, we created the frontend only with React
- Then we created an OpenAPI specification for defining the frontend-backend API
- Next, we created the backend from this specification
- We added database support with SQLite and SQLAlchemy

You can find the code in the [interview-canvas-share](#) repository.

Containerization

When we run the application locally, we need to execute two commands: one for the frontend and one for the backend.

Let's run them in two separate terminals:

```
# first terminal
cd frontend
npm run dev


# second terminal
make run
```

So, we have two services, and we need to deploy them to production. We may think that we need two containers: one for the frontend and one for the backend.

During development, we run the frontend as a separate service because we use Vite. Vite watches the frontend code and refreshes the page when we make changes. It's very convenient because we can see our changes immediately.

In production, the frontend code doesn't change while the application is running. We build it once and get a set of static HTML, CSS, and JavaScript files.

This means we don't need a separate container for the frontend: the backend can serve these files. So, we need only one container.



In development (1), the frontend and backend run separately. In production (2, 3), FastAPI serves the built frontend from one container

Ask the coding assistant to create it:

Create a Dockerfile that builds the frontend with Node, then builds a Python image with the backend and the frontend static files.

Backend should serve the frontend.

We get a two-stage Docker build:

- First, we use a Node.js image to compile the frontend
- Then we build the backend and copy only the frontend files without the Node.js dependencies

You can see the [Dockerfile here](#).

Build and run the image from the repository root:

```
docker build -t sdip:latest .

docker run --rm -p 8000:8000 \
  -v sdip-data:/data \
  -e SDIP_DATABASE_URL=sqlite:///data/sdip.db \
  --name sdip sdip:latest
```

Here we specify a named Docker volume `sdip-data` that will keep the SQLite database between container runs.

Open the application at localhost:8000 and test it:

- Create an interview session
- Open the join link in a different browser window
- Move an element in the candidate window
- Check that the interviewer sees the change

We'll repeat this test again. I'll refer to it as the two-session test.

Switch from SQLite to Postgres

SQLite is very convenient for local development. It keeps the data in a single file and doesn't need a separate database server.

But for production, we typically use Postgres or a similar database.

When we set up the foundation [in the previous article](#), we asked the coding agent to use SQLAlchemy.

I did this on purpose because I knew that later I'd switch to Postgres.

Start Postgres locally:

```
docker run -d \  
  --name interview-canvas-db \  
  -e POSTGRES_USER=sdip \  
  -e POSTGRES_PASSWORD=sdip \  
  -e POSTGRES_DB=sdip \  
  -p 5432:5432 \  
  -v interview-canvas-pgdata:/var/lib/postgresql/data \  
  postgres:16-alpine
```

Now we can ask the assistant to use it:

```
Add Postgres support to the backend.
```

Run the backend against Postgres from the repository root:

```
export SDIP_DATABASE_URL=postgresql://sdip:sdip@localhost:5432/sdip  
make run
```

When it's done, repeat the two-session test.

Docker Compose

Previously, I started a Postgres container with a separate command. But now let's put all the services our application needs inside one Docker Compose file.

With this file, we can run our entire application with a single command `docker compose up`.

Ask the assistant to implement it:


```
Create docker-compose.yaml with two services: Postgres and the app.
```

The file defines the database and adds a health check, so our application waits until Postgres is ready to accept connections.

Start it:

```
docker compose up --build
```

In our case, it runs the application at localhost:8100.



We run the application with Postgres using Docker Compose

Integration and end-to-end tests

The AI assistant might have created some backend tests.


If it didn't, ask it to create integration tests:

```
Create integration tests that run against docker-compose.yaml. What scenarios should we test?
```

We added two things that can potentially break, so let's test them too. We'll verify that:

- The frontend compiles correctly

- The backend can communicate with Postgres



Playwright can run the two-session test for us

Ask the AI assistant to implement a test:

Add an end-to-end test that runs against `docker-compose.yaml`.

Use Playwright to:

1. Log in as the interviewer (session 1).
2. Create an interview session.
3. Share the join link.
4. Join from a separate client as the candidate (session 2).
5. Change the canvas as the candidate (session 2).
6. Verify that the interviewer sees the change (session 1).

Put the tests in the `e2e/` folder in the repository root.

After it finishes, we can run the tests with a single make command:

```
make e2e
```

Deploy to AWS

We're now certain that the application works well. We can deploy it.


Our application runs in a container and only needs Postgres, so we have a lot of options for deploying it. We can use Render, Railway, Fly.io, or any other managed container system.

[Last year we deployed to Render](#), but this year I want to deploy to AWS. You don't have to use AWS, and instead you can ask the coding assistant to recommend an environment for your application.

Ask your assistant to deploy it:

```
Deploy this application to AWS. Use AWS CloudFormation.
```

For that to work, you need to have an AWS user. I typically create a temporary user with admin permissions and watch every step of what the coding agents are doing.



One EC2 instance runs Caddy, the app, and Postgres. We manage it through CloudFormation.

You can see what I got [here](#). It runs the app, Postgres, and [Caddy](#) (adds HTTPS and WSS support for our app) on one EC2 instance.

It's fine for a proof-of-concept, but using managed database services (such as RDS) is better. We will not do it here.

CI/CD with GitHub Actions

We used a user with the admin permissions to deploy the application. It's okay for the first deployment, but only when we carefully watch it. The next step is to configure CI/CD and remove that access.

- Continuous integration (CI) means that every time we make a change and push it to GitHub, we automatically run all the tests to make sure we didn't break anything.
- Continuous deployment (CD) is about deploying this change automatically.


In GitHub, we use GitHub Actions for that.

Let's configure it. Every time we make a push to main, we want to:

- run frontend and backend tests

- build the containers
- run the integration tests
- run the end-to-end tests
- if all the tests pass, deploy the new version

For the last step, the runner (the process that will deploy the application) will need to be able to access our AWS infrastructure. We will use OpenID Connect (OIDC) for this: the runner will assume a role with the necessary permissions and update the application.



GitHub Actions runs frontend and backend tests in parallel, checks the full stack, and then deploys to AWS

Create the role and the workflow:

Create a CI/CD pipeline that:

- runs backend and frontend tests in parallel
- builds the Docker Compose stack and runs integration and end-to-end tests against it
- deploys to AWS using a GitHub OIDC role
- validates that the deploy is successful by checking the health endpoint

The [finished workflow](#) uses a restricted AWS role for deployment.

Change something in the application, commit, and push to see it go live.

Clean up


When we're done, we need to delete all the resources created by CloudFormation:

```
aws cloudformation delete-stack --stack-name sdip
aws cloudformation wait stack-delete-complete --stack-name sdip
```

The next step after deployment

Let's recap what we have done so far:


- We created the frontend application with React
- Then we defined the frontend-backend API with OpenAPI specs
- Based on the specs, we created the backend
- Next, we added database support with SQLite and SQLAlchemy to the backend
- To make it easier to deploy the application, we put both frontend and backend inside one container. The backend serves the frontend.
- To go to production, we replaced SQLite with Postgres.
- Next, we simplified running everything locally with Docker Compose
- Once everything was in Compose, we created an end-to-end test
- We took the application that worked locally and deployed it to AWS using CloudFormation
- Finally, we created a CI/CD deployment pipeline to deploy every change automatically.



The end-state architecture: two browser sessions connect through Caddy to the app and persistent Postgres data

But we still need to do a lot more for the app to run reliably:

- Dev and prod environments
- Observability: logs, metrics, alerts
- Using AI as the first responder when the application stops working



The user promotes releases while separate observe and respond loops detect production problems and create fixes

We'll cover it in the next lesson.



DevOps and Observability for an AI-Built App

ALEXEY GRIGOREV • 18 AOÛT

[Read full story →](#)

Thanks for reading Alexey On Data! Subscribe for free to receive new posts and support my work.



19 Likes • 3 Restacks

[← Previous](#)

[Next →](#)

Discussion about this post

[Comments](#) [Restacks](#)



Write a comment...

